

# An Investigation into Parareal

Abhijit Chowdhary

New York University

May 2019

# Problem Statement

We would like to numerically solve the ordinary differential equation:

$$\begin{cases} u'(t) = f(t, u), & t \in [t_0, t_f] \\ u(t_0) = u_0 \end{cases}$$

where  $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$  and  $u : \mathbb{R} \rightarrow \mathbb{R}^d$ .

Framing this in a HPC context, how can we solve the above while taking advantage of massively parallel hardware?

# Serial Methods

- Forward and Backward Euler:

$$u_{n+1} = u_n + hf(t, u_n), u_n + hf(t, u_{n+1})$$

- Linear Multistep Methods

$$\sum_k \alpha_k y_{n+k} = h \sum_k \beta_k f(t_{n+k}, y_{n+k})$$

- Runge Kutta Methods
- Etc.

# Parallel Techniques

## Problem!

Most of the previous methods are iterative, with dependency on the previously computed values.

Not embarrassingly parallel... We have to get creative.

# Parallel-In-Time

One technique is to parallelize the problem across time, i.e. split  $[t_0, t_f]$  into slices  $[t_n, t_{n+1}]$  and in parallel solve each slice.

## Iteration Dependence

But since each slice depends on the previous, how can we chain together these individual solutions to each slice?

## Predictor-Corrector

What if we were to predict the portions that  $[t_n, t_{n+1}]$  are dependent on with a cheap coarse operator, and then correct with a fine operator!

Yes!

# Parareal


Let  $\mathcal{G}$  be a course and inexpensive operator, and let  $\mathcal{F}$  be a fine operator of high order. Then Parareal iteration is:

$$\lambda_{n+1}^{k+1} = \mathcal{G}(t_{n+1}, t_n, \lambda_n^{k+1}) + \left[ \mathcal{F}(t_{n+1}, t_n, \lambda_n^k) - \mathcal{G}(t_{n+1}, t_n, \lambda_n^k) \right]$$

## Parallel Potential

Notice! The  $\mathcal{F}$  term depends only on the previous solutions, which means it can be computed in parallel for each  $k$  step. Furthermore,  $\mathcal{G}(\lambda_n^k)$  satisfies the FSAL property!


# Parareal: Visual Example



Full beautiful animation found on wikipedia's Parareal page.




# Parareal: Visual Example



Full beautiful animation found on wikipedia's Parareal page.

# Parareal: Visual Example



Full beautiful animation found on wikipedia's Parareal page.

# Pseudocode

$$\lambda_{n+1}^{k+1} = \mathcal{G}(t_{n+1}, t_n, \lambda_n^{k+1}) + [\mathcal{F}(t_{n+1}, t_n, \lambda_n^k) - \mathcal{G}(t_{n+1}, t_n, \lambda_n^k)]$$

**Require:**  $y_0$  and course and fine solvers  $\mathcal{G}$ ,  $\mathcal{F}$ .

$y_c \leftarrow \mathcal{G}(t_f, t_0, y_0).$  ▷ Coarsely approximate solution

$y \leftarrow y_c.$

**while** iter < max\_iter && not converged **do**

**for**  $n = 0 \rightarrow P$  **do** ▷ Parallel capable

$y_f(n) = \mathcal{F}(t_{n+1}, t_n, y(n)).$

$\delta y(n) = y_f(n) - y_c(n).$  ▷ corrector term. FSAL

**end for**

**for**  $n = 0 \rightarrow P$  **do**

$y_c(n) = \mathcal{G}(t_{n+1}, t_n, y(n)).$  ▷ Predict.

$y(n) = y_c(n) + \delta y(n).$  ▷ Correct.

**end for**

**end while**

# Parallel Pseudocode

**Require:**  $y_0$  and coarse and fine solvers  $\mathcal{G}$ ,  $\mathcal{F}$ .

$y_c \leftarrow \mathcal{G}(t_f, t_0, y_0).$  ▷ Coarsely approximate solution

$y \leftarrow y_c.$

**while** iter < max\_iter && not converged **do**

  #pragma omp parallel for

**for**  $n = 0 \rightarrow P$  **do**

$y_f(n) = \mathcal{F}(t_{n+1}, t_n, y(n)).$

$\delta y(n) = y_f(n) - y_c(n).$  ▷ corrector term. FSAL

**end for**

**for**  $n = 0 \rightarrow P$  **do**

$y_c(n) = \mathcal{G}(t_{n+1}, t_n, y(n)).$  ▷ Predict.

$y(n) = y_c(n) + \delta y(n).$  ▷ Correct.

**end for**

**end while**

# Speedup Analysis

Suppose we have  $P$  processors, and suppose our fine method takes  $T_f$  time,  $T_g$  for the course method. Furthermore, assume that our Parareal iteration needs  $k$  steps. Then the speedup it provides is:

$$S = \frac{PT_f}{PT_g + k(PT_g + T_f)} = \frac{1}{\frac{T_g}{T_f} + k(\frac{T_g}{T_f} + \frac{1}{P})} = \frac{1}{\frac{T_g}{T_f}(1+k) + \frac{k}{P}}$$

Note:

$$S_\infty = \lim_{P \rightarrow \infty} \frac{1}{\frac{T_g}{T_f}(1+k) + \frac{k}{P}} = \frac{T_f}{T_g(1+k)}$$

# No Free Lunch

$$S_{\infty} = \frac{T_f}{T_g(1+k)}$$


This tells us something very important, we want the ratio  $T_f/T_g$  to be as large as possible, and we want  $k$  to be as small as possible.

## Theorem


The parareal method has order of accuracy  $mk$ , where  $k - 1$  is the number of parareal iterations made and  $m$  is the order of  $\mathcal{G}$ , assuming  $\mathcal{F}$  is close to truth. (Bal)

Therefore, we have this uncomfortable optimization problem. We want  $k$  small for speedup, but large for order. If we make  $m$  large instead, then the ratio  $T_f/T_g$  becomes smaller...

# Forward Euler Convergence



# Runge-Kutta Convergence





# Improvements | Pipelined Parareal




Figure borrowed from Ruprecht's *Implementing Parareal*.


# Pipelined Parareal Pseudocode

It's a bit too complicated to present in a slide, but the general idea is to begin the omp parallel statement at the beginning and then pretend like were in a MPI environment.


## General Idea

Thread  $p$  processes the new iteration at  $p + 1$ . MPI communication is mimicked with locks, and done in place.

# Strong Scaling



# Weak Scaling



# Conclusions

- We see significant speedup as well as proper weak and strong scaling of parareal, supposing ideal choices of  $\mathcal{F}, \mathcal{G}$ .

# Conclusions

- We see significant speedup as well as proper weak and strong scaling of parareal, supposing ideal choices of  $\mathcal{F}, \mathcal{G}$ .
- The algorithm is very picky with choices of  $\mathcal{G}$  and  $\mathcal{F}$ , and choices of  $\Delta t$  and  $\delta t$  vary wildly from problem to problem.

# Conclusions

- We see significant speedup as well as proper weak and strong scaling of parareal, supposing ideal choices of  $\mathcal{F}, \mathcal{G}$ .
- The algorithm is very picky with choices of  $\mathcal{G}$  and  $\mathcal{F}$ , and choices of  $\Delta t$  and  $\delta t$  vary wildly from problem to problem.
- To generalize to a larger scale, the code should be formulated first with MPI to distribute work across nodes, and then with OpenMP to parallelize locally.